


智能体记忆的研究-RAG 核心

1. 概览

智能体的 **智能** 的核心影响因素是什么？



2. 记忆原理和机制





Figure 1: A unified framework of memory Taxonomy, Operations, and Applications in AI systems.

2.1 记忆/存储(memory)表示分类



2.2 六大核心操作

仅仅把信息存起来还不够，AI 还需要对这些记忆进行有效的管理和利用。论文借鉴了认知心理学的思想，将 AI 的记忆操作分为了两大类共六种：

2.2.1 记忆管理 (Memory Management)

这类操作主要负责记忆的存储、维护和清理。提供四种操作（巩固、索引、更新、遗忘）整理成清晰结构的表格，便于理解其定义、作用及应用场景：

操作名称	定义与说明	作用与意义	示例
巩固 (Consolidation)	将短期信息转化为更持久、稳定的记忆形式。可能涉及将对话要点存入长期知识库或将新知识融入模型参数。	实现长期学习、个性化和构建持久知识体系的关键步骤。	学习助手在多次互动后总结学生的薄弱知识点，并记录到「学习档案」中作为长期记忆。
索引 (Indexing)	为存储的记忆内容创建目录或标签，便于快速、准确地检索。可以基于关键词、实体、时间戳、关系等。	提高记忆检索效率和准确性，尤其适用于大规模记忆库。	客服机器人根据用户 ID、对话时间、产品类型等信息建立索引，加快问题查找和处理速度。
更新 (Updating)	当新信息出现或原有记忆有误时，对已有记忆进行修改、补充或替换。如模型编辑或数据库条目更新。	保持记忆的准确性与时效性，使 AI 能适应环境变化和知识演进。	AI 发现某公司 CEO 更换后，将原来的「张三」更新为「李四」。
遗忘 (Forgetting)	选择性移除或抑制过时、不相关、错误或需删除的记忆内容（如出于隐私考虑）。	避免信息过载、减少干扰、保护隐私，提升记忆系统的整体性能。	用户要求删除聊天记录；新闻推荐 AI 忽略用户很久未关注的话题。

2.2.2 记忆利用 (Memory Utilization)

这类操作主要负责在需要时提取和使用已存储的记忆。

以下是两个操作（**检索**和**压缩**）与之前内容整合后的完整表格，涵盖了 AI 记忆管理系统的大大核心操作清晰展示每个操作的功能、意义与示例：

操作名称	定义与说明	作用与意义	示例
检索 (Retrieval)	当 AI 面临一个任务或查询时，从其记忆库中识别并提取出最相关的记忆内容。	是 AI 利用记忆辅助思考、决策和生成回应的前提。	用户问：「我上次跟你聊到我喜欢吃什么？」AI 回顾对话历史找到答案。
压缩 (Compression)	在保留核心信息的前提下，减少记忆内容的体积。可通过对文本摘要、关键信息提取等方式实现。	解决上下文窗口限制，提高处理效率，节省资源成本。	AI 阅读长篇报告后生成简短摘要，或将长对话历史压缩为关键点后再送入模型处理。

2.2.2.1 检索 (Retrieval)

$$\begin{aligned} \text{Retrieve}(\mathcal{M}_t, \mathcal{Q}) &= m_{\mathcal{Q}} \in \mathcal{M}_t \\ &\text{with } \text{sim}(\mathcal{Q}, m_{\mathcal{Q}}) \geq \tau \end{aligned} \tag{5}$$

Yiming Du, Hongru Wang, Zhengyi Zhao, Bin Liang, Baojun Wang, Wanjun Zhong, Zezhong Wang, and Kam-Fai Wong. 2024. PerlTqa: A personal long-term memory dataset for memory classification, retrieval, and synthesis in question answering. *arXiv preprint arXiv:2402.16288*.

Bingbing Wang, Yiming Du, Bin Liang, Zhixin Bai, Min Yang, Baojun Wang, Kam-Fai Wong, and Ruifeng Xu. 2025a. A new formula for sticker retrieval: Reply with stickers in multi-modal and multi-session conversation. In *Proceedings of the AAI Conference on Artificial Intelligence*, volume 39, pages 25327– 25335.

Junjie Zhou, Zheng Liu, Shitao Xiao, Bo Zhao, and Yongping Xiong. 2024. Vista: Visualized text embedding for universal multi-modal retrieval. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3185–3200.

Hexiang Tan, Fei Sun, Wanli Yang, Yuanzhuo Wang, Qi Cao, and Xueqi Cheng. 2024b. Blinded by generated contexts: How language models merge generated and retrieved contexts when knowledge conflicts? In *Proceedings of the 62nd Annual Meeting of the Association for*

Computational Linguistics (Volume 1: Long Papers), pages 6207–6227, Bangkok, Thailand. Association for Computational Linguistics.

Bingbing Wang, Yiming Du, Bin Liang, Zhixin Bai, Min Yang, Baojun Wang, Kam-Fai Wong, and Ruifeng Xu. 2025a. A new formula for sticker retrieval: Reply with stickers in multi-modal and multi-session conversation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 25327–25335.

2.2.2.2 压缩 (Compression)

压缩通过在将信息输入模型之前以压缩率 α 保留显著信息并丢弃冗余信息，从而实现在有限上下文窗口下高效利用上下文。压缩大致可分为输入前压缩和检索后压缩。输入前压缩适用于无需检索的长上下文模型，其中对全上下文输入进行评分、过滤或汇总，以符合上下文约束。检索后压缩在记忆访问之后进行，通过模型推理前的上下文压缩或通过将检索到的知识整合到模型参数中来减少检索到的内容。与在记忆构建过程中总结信息的记忆巩固不同，压缩侧重于在推理过程中减少记忆。

$$\mathcal{M}_t^{comp} = \text{Compress}(\mathcal{M}_t, \alpha)$$

Haofei Yu, Cunxiang Wang, Yue Zhang, and Wei Bi. 2023. TRAMS: Training-free memory selection for long-range language modeling. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 4966–4972, Singapore. Association for Computational Linguistics.

Tsz Ting Chung, Leyang Cui, Lemao Liu, Xinting Huang, Shuming Shi, and Dit-Yan Yeung. 2024. Selection-p: Self-supervised task-agnostic prompt compression for faithfulness and transferability. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 11057–11070, Miami, Florida, USA. Association for Computational Linguistics.

Fangyuan Xu, Weijia Shi, and Eunsol Choi. 2024a. RE-COMP: Improving retrieval-augmented LMs with context compression and selective augmentation. In *The Twelfth International Conference on Learning Representations*.

Ali Safaya and Deniz Yuret. 2024. Neurocache: Efficient vector retrieval for long-range language modeling. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 870–883, Mexico City, Mexico. Association for Computational Linguistics.

Wanjun Zhong, Lianghong Guo, Qiqi Gao, He Ye, and Yanlin Wang. 2024. Memorybank: Enhancing large language models with long-term memory. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 19724–19731.

Kuang-Huei Lee, Xinyun Chen, Hiroki Furuta, John Canny, and Ian Fischer. 2024b. A human-inspired reading agent with gist memory of very long contexts. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 26396–26415. PMLR.

2.3 主要的应用场景

2.3.1 长期记忆 (Long-Term Memory)

长期记忆针对问题的是：AI 系统如何存储、管理和利用跨越多次交互、甚至很长时间跨度的信息。目标是让 AI 像人一样，能够记住过去的经验，并用于指导未来的行为，实现真正的个性化和持续学习。

2.3.1.1 llm-powered personalized agent

开放域对话系统随大语言模型发展取得显著进步，但多数现有系统主要关注简短单会话交互，忽视了与聊天机器人长期陪伴和个性化交互的现实需求，而事件总结和角色管理对满足这一需求至关重要。鉴于 LLM 类人认知和推理能力的最新进展，研究人员引入了与模型无关的框架“长期对话代理 (LD-Agent)”，其包含事件感知、角色提取和响应生成三个独立可调模块，事件记忆模块采用长短期记忆库分别关注历史和进行中的会话，并引入基于主题的检索机制提高记忆检索准确性，角色模块对用户和代理进行动态角色建模，检索的记忆和提取的角色集成后输入生成器以生成适当响应，实验证明了 LD-Agent 在各种基准、模型和任务中的有效性、通用性和跨域能力，相关代码已发布。




Figure 2: The Framework of LD-Agent. The event module stores historical memories from past sessions in long-term memory and current context in short-term memory. The persona module dynamically extracts and updates personas for both users and agents from ongoing utterances, storing them in a persona bank for each character. The response module then synthesizes this data to generate informed and appropriate responses.

Hao Li, Chenghao Yang, An Zhang, Yang Deng, Xiang Wang, and Tat-Seng Chua. 2024a. Hello again! llm-powered personalized agent for long-term dialogue. *arXiv preprint arXiv:2406.05925*. <https://arxiv.org/abs/2406.05925>

Di Wu, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu. 2024a. Longmemeval: Benchmarking chat assistants on long-term interactive memory. *arXiv preprint arXiv:2410.10813*.

Wanjun Zhong, Lianghong Guo, Qiqi Gao, He Ye, and Yanlin Wang. 2024. Memorybank: Enhancing large language models with long-term memory. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 19724–19731.

Jing Xu, Arthur Szlam, and Jason Weston. 2021. Beyond goldfish memory: Long-term open-domain conversation. *arXiv preprint arXiv:2107.07567*. <https://arxiv.org/abs/2107.07567>

Adyasha Maharana, Dong-Ho Lee, Sergey Tulyakov, Mohit Bansal, Francesco Barbieri, and Yuwei Fang. 2024. Evaluating very long-term conversational memory of llm agents. *arXiv preprint arXiv:2402.17753*.

2.3.2 长上下文记忆 (Long-Context Memory)

长上下文记忆用于 AI 系统处理和理解非常长的输入序列，比如整本书、万字长文、数小时的会议记录等。主要关注上下文非结构化记忆的处理，也与参数化记忆的利用效率（如 Transformer 中的 KV 缓存）有关。核心操作是压缩与检索。

2.3.2.1 H2o: Heavy-hitter oracle for efficient generative inference of large language models

大型语言模型（LLMs）部署成本高昂，尤其是长内容生成场景，其 KV 缓存会随序列长度和批量大小线性占用 GPU 内存。本文提出 H₂O 方法，基于“少量 token 在计算注意力分数时贡献主要价值（即 Heavy Hitters, H₂）”的观察，设计动态保留近期和 H₂ token 的 KV 缓存逐出策略，将其建模为动态子模问题并给出理论保证。在 OPT、LLaMA、GPT - NeoX 等模型上验证显示，H₂O 用 20% 的 H₂ token，使 OPT - 6.7B 和 OPT - 30B 在 DeepSpeed Zero - Inference、Hugging Face Accelerate、FlexGen 等推理系统上的吞吐量提升最高达 29×、29×、3×，同批量下延迟降低 1.9×，代码已开源。




Figure 1: Upper plots illustrate symbolic plots of an attention map deploying different KV cache policies in LLM generation. Lower right: contrasts their accuracy-memory trade-off. Left: the overview of H₂O framework.

Zhenyu Zhang, Ying Sheng, Tianyi Zhou, Tianlong Chen, Lianmin Zheng, Ruisi Cai, Zhao Song, Yuan-dong Tian, Christopher Re, Clark Barrett, Zhangyang Wang, and Beidi Chen. 2023b. H2o: Heavy-hitter oracle for efficient generative inference of large language models. In *Thirty-seventh Conference on Neural Information Processing Systems*.

2.3.2.2 XRAG: Extreme context compression

xRAG 是专为检索增强生成 (Retrieval-Augmented Generation, RAG) 设计的上下文压缩技术。传统 RAG 通过检索文档并拼接文本作为上下文输入语言模型，而 xRAG 的目标是通过压缩上下文提升效率，同时保持生成质量。




Figure 2: Overview of xRAG (a) and RAG (b). For a given query, RAG typically concatenates the retrieved document with the query, significantly extending the context length. In contrast, xRAG addresses this issue through modality fusion by directly projecting the document embedding into the LLM’s representation space. This allows for efficient retrieval-augmentation with the addition of only one token.

- xRAG 创新：将嵌入视为一种独立的检索模态特征（类似图像、音频等模态），而非仅作为检索工具。这一重新定义打破了“嵌入→文本”的依赖关系，为压缩上下文提供了可能。
- 技术路径：通过模态融合（Modality Fusion）将检索模态的嵌入直接映射到语言模型的语义空间，使嵌入特征能被语言模型理解和利用。
- 可训练组件：仅训练一个模态桥（Modality Bridge），用于实现检索嵌入到语言模型空间的映射，参数规模极小。
- 优势：
 - 可复用离线构建的文档嵌入库，无需在线计算嵌入，降低部署成本；
 - 保持 RAG 的“即插即用”特性，无需修改原有检索系统架构。

Xin Cheng, Xun Wang, Xingxing Zhang, Tao Ge, Si- Qing Chen, Furu Wei, Huishuai Zhang, and Dongyan Zhao. 2024. xrag: Extreme context compression for retrieval-augmented generation with one token. *arXiv preprint arXiv:2405.13792*. <https://arxiv.org/abs/2405.13792>

Huiqiang Jiang, Qianhui Wu, Xufang Luo, Dong- sheng Li, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. 2024a. LongLLMLingua: Accelerating and enhancing LLMs in long context scenarios via prompt compression. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1658–1677, Bangkok, Thailand. Association for Computational Linguistics.

2.3.3 参数化记忆修改 (Parametric Memory Modification)

直接对 AI 模型（尤其是 LLM）内部参数中存储的知识进行修改。这就像对 AI 的「大脑」进行外科手术。核心记忆类型是参数化记忆。主要操作涉及更新、遗忘和巩固。

2.3.3.1 通过 Null Space 修改参数输出

大型语言模型（LLMs）常因知识错误或过时出现幻觉问题，模型编辑方法应运而生，其中主流的“先定位后编辑”范式虽有效，但当前研究表明其引入的扰动会破坏模型原有知识，尤其在顺序编辑场景中。为此，AlphaEdit 提出将扰动投影到保留知识的零空间后再应用于参数，从理论上确保编辑后模型在查询保留知识时输出不变，缓解知识破坏问题。在 LLaMA3、GPT2 - XL 和 GPT - J 等多种 LLMs 上的实验显示，AlphaEdit 仅通过一行额外的投影代码，就能将大多数“先定位后编辑”方法的性能平均提升 36.7%，其代码已开源。




Figure 3: Comparison between the paradigms of AlphaEdit and current method. Best viewed in color.


Junfeng Fang, Houcheng Jiang, Kun Wang, Yunshan Ma, Jie Shi, Xiang Wang, Xiangnan He, and Tat-Seng Chua. 2025. Alphaedit: Null-space constrained model editing for language models. In *The Thirteenth International Conference on Learning Representations*.

<https://arxiv.org/abs/2410.02355>

2.3.3.2 MASS-EDITING MEMORY IN A TRANSFORMER

近期研究在利用新记忆更新大型语言模型以替换过时信息或添加专业知识方面展现出令人兴奋的前景，但该领域工作主要局限于更新单一关联。为此，研究者开发了 MEMIT 方法，用于直接用许多记忆更新语言模型，实验表明该方法可扩展到 GPT-J（6B）和 GPT-NeoX（20B）的数千个关联，较之前的工作有数量级的超越，其代码和数据可在指定链接获取。

主要突破是将大模型记忆更新的规模从单一关联扩展到数千个关联，实验验证了方法在 60 亿参数（GPT-J）和 200 亿参数（GPT-NeoX）模型上的有效性，为大模型知识大规模更新提供了可行方案。



Kevin Meng, Arnab Sen Sharma, Alex Andonian, Yonatan Belinkov, and David Bau. 2022b. Mass-editing memory in a transformer. *arXiv preprint arXiv:2210.07229*.

Pratyush Maini, Zhili Feng, Avi Schwarzschild, Zachary Chase Lipton, and J Zico Kolter. 2024. TOFU: A task of fictitious unlearning for LLMs. In *First Conference on Language Modeling*.


Yu Wang, Chi Han, Tongtong Wu, Xiaoxin He, Wangchunshu Zhou, Nafis Sadeq, Xiusi Chen, Zexue He, Wei Wang, Gholamreza Haffari, et al. 2024j. Towards lifespan cognitive systems. *arXiv preprint arXiv:2409.13265*.

2.3.4 多源记忆 (Multi-Source Memory)

多源记忆主要针对的问题是，AI 系统如何整合和利用来自不同来源、不同类型（甚至不同模态）的记忆信息进行推理和决策。涉及的记忆类型主要是上下文结构化记忆（如知识图谱、数据库）和上下文非结构化记忆（如网页文本、用户对话），有时也需要结合模型自身的参数化记忆（先验知识）。核心记忆操作包括检索、更新和巩固。

2.3.4.1 Chatdb: Augmenting llms with databases as their symbolic memory.


具备记忆的大型语言模型（LLMs）在计算上具有通用性，但主流 LLMs 未能充分利用记忆，且其设计深受生物大脑影响，传统神经记忆机制因具有近似性且易累积错误，无法支持 LLMs 进行复杂推理。为此，论文从现代计算机架构中获取灵感，提出通过将 SQL 数据库作为符号记忆来增强 LLMs，使 LLM 生成 SQL 指令以操作数据库，从而实现复杂多跳推理，并在需要复杂推理的合成数据集上验证了该记忆框架的有效性，其项目网站也已公开。



Chenxu Hu, Jie Fu, Chenzhuang Du, Simian Luo, Junbo Zhao, and Hang Zhao. 2023. Chatdb: Augmenting llms with databases as their symbolic memory.

2.3.4.2 A new formula for sticker retrieval

这篇论文介绍了面向多模态多会话场景的贴纸检索研究，作者创建了包含 1542 个会话、50192 条话语和 2182 个贴纸的新数据集 MultiChat，并提出 Intent-Guided Sticker Retrieval (IGSR) 框架，该框架通过引入贴纸属性构建记忆库，结合当前会话提取相关记忆推导意图，进而检索贴纸，实验表明 IGSR 在 MultiChat 数据集上性能优于基线模型，验证了其在多模态多会话场景的有效性



Bingbing Wang, Yiming Du, Bin Liang, Zhixin Bai, Min Yang, Baojun Wang, Kam-Fai Wong, and Ruifeng Xu. 2025a. A new formula for sticker retrieval: Reply with stickers in multi-modal and

multi-session conversation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 25327– 25335.

3. 人体的记忆机制

人体的记忆是一个复杂的心理过程，它允许我们存储和回忆信息。记忆可以按照不同的标准进行分类，并且其运作原理涉及大脑多个区域的协同工作。

3.1 记忆的分类

1. 根据时间长短

- **感觉记忆（瞬时记忆）**：持续时间极短，通常不到一秒到几秒，是对感官刺激的即时印象。
- **短期记忆（工作记忆）**：能够保持几秒至几分钟的信息，容量有限，通常能记住 7 ± 2 个信息块。
- **长期记忆**：可以保存信息从数分钟到终身不等。它的容量几乎是无限的。

2. 根据内容类型

- **陈述性记忆（显式记忆）**：关于事实和事件的记忆，又分为语义记忆（例如知道巴黎是法国的首都）和情景记忆（个人经历的具体事件）。
- **非陈述性记忆（隐式记忆）**：包括技能和习惯、条件反射等，不需要有意识地回想即可表现出来，比如骑自行车或打字。

3. 根据处理方式

- **程序性记忆**：与如何做某事相关的记忆，如驾驶汽车。
- **情感记忆**：与特定情绪状态相关联的记忆。
- **形象记忆**：基于视觉或听觉等感官输入的形象化记忆。

3.2 记忆的工作原理

生物记忆的原理涉及大脑如何通过改变神经元之间的连接来存储信息。这个过程非常复杂，涉及多种机制和分子层面的变化。以下是一些关键概念和技术细节，有助于理解生物记忆的工作原理：

分类	描述	详细说明
神经可塑性		

	大脑改变其结构和功能的能力，是学习和记忆的基础。	- 突触可塑性 ：突触强度的变化，包括长期增强（LTP）和长期抑制（LTD）。- 结构可塑性 ：神经元之间形成新连接或修剪旧连接。
长期增强作用（LTP）	突触水平的重要记忆机制。	发生在频繁同步激活的神经元之间，通过改变突触后受体密度和敏感度来增强连接强度。特别涉及NMDA和AMPA受体，当突触持续激活时，引发一系列生化反应，导致更多AMPA受体插入到突触后膜中，从而增强信号传导。
蛋白质合成与基因表达	维持LTP并转化为长期记忆需要新的蛋白质合成和特定基因表达。	涉及早期响应基因（如c-Fos）的激活，这些基因编码转录因子和其他蛋白质，进一步调控其他基因的表达，最终支持记忆的巩固过程。
海马体的作用	在新记忆形成中扮演重要角色，特别是情景记忆和空间记忆。	海马体损伤影响新长期记忆的形成，但不影响短期记忆或已有的长期记忆。参与短时记忆向长时记忆转化的过程，即“记忆巩固”。
记忆巩固	学习体验到稳定存储于大脑中的过程。	初期依赖海马体的支持，随着时间推移，记忆逐渐转移到皮层区域进行更持久的存储。这一过程可能需要几天甚至数周，并且涉及多次重复回忆和再巩固。

3.3 记忆的生命周期

我们可以将记忆的完整生命周期扩展为以下五个阶段：

阶段	定义与说明	关键特征与影响因素
编码（Encoding）	将外界信息转化为大脑可以处理的形式。	涉及注意力、感知和初步神经活动模式的形成。
巩固（Consolidation）	将短期记忆转化为更稳定的长期记忆。	主要在睡眠中发生（如慢波睡眠SWS、REM）；海马体与皮层交流至关重要。
存储（Storage）	长期记忆被分散存储在大脑的不同区域。	不同类型的记忆分布在不同脑区（如海马体、大脑皮层、小脑等）。
检索（Retrieval）	当需要使用信息时，从存储中找回记忆。	依赖于编码质量、提示线索、当前情绪和环境状态。
遗忘（Forgetting）	信息无法被检索或逐渐减弱的过程。	并非总是坏事，是大脑选择性优化、避免信息过载的一种机制。

3.3.1 遗忘的类型与机制

遗忘类型	定义说明	子类或机制	相关脑区/物质
主动遗忘（Active Forgetting）	大脑有意识或无意识地抑制某些记忆	- 抑制机制：压制不相关或干扰记忆- 更新机制：新信息替代旧信息	前额叶、多巴胺、GABA
被动遗忘（Passive Forgetting）	由于缺乏复习或激活，记忆痕迹逐渐衰退	- “用进废退”原则- 痕迹衰退理论	—
干扰遗忘（Interference Theory）	新旧记忆之间相互干扰导致提取失败	- 前摄干扰：旧记忆干扰新学习- 倒摄干扰：新记忆干扰旧回忆	海马体等记忆相关脑区
提取失败（Retrieval Failure）	记忆仍然存在，但因线索不足或情境变化无法提取	- 缺乏合适的提示- 情境不匹配（如“舌尖现象” Tip-of-the-Tongue, TOT)	—
再巩固过程中的修改或削弱	回忆记忆时其进入不稳定状态，需重新巩固	- 可能被修改、强化、削弱或部分丢失	海马体、神经递质、蛋白质合成相关机制

3.3.2 遗忘的生物学基础

- **突触可塑性变化**：未使用的突触连接会被修剪。
- **蛋白质降解**：维持记忆所需的蛋白质若未持续合成，可能导致记忆消失。
- **神经网络重组**：随着大脑结构的变化，一些旧的神经通路可能不再活跃。
- **海马体和皮层交互变化**：长期记忆的迁移过程中，若未能有效整合，也可能导致遗忘。

遗忘的意义

正面意义	解释
清理冗余信息	减少认知负担，提高效率
避免错误记忆	抑制冲突或混淆的记忆
促进决策灵活性	忘记过去的选择有助于适应新环境
心理健康保护	抑制创伤记忆有助于缓解心理压力

2.4. 个体智商的相关因素

一个生物个体的智商（IQ）与其大脑的多个方面有关，主要包括以下几点：

- 1. **脑容量：**虽然脑容量与智能之间并非直接的线性关系，但通常来说，较大的脑容量可能意味着更多的神经元和更复杂的神经网络，这为更高的认知功能提供了物理基础。不过，重要的是脑结构的复杂性和效率，而不仅仅是大小。
- 2. **大脑皮层：**特别是大脑皮层（尤其是额叶），在高级认知功能中扮演关键角色，包括决策、问题解决和社会行为等。大脑皮层的厚度和褶皱程度（即表面积）往往与智能水平有正相关的关系。
- 3.Disclaimer: The original text contained an incomplete statement regarding the relationship between neuronal connections and intelligence, which has been omitted for clarity.
- 4. **神经可塑性：**这是指神经系统通过改变其结构和功能以适应环境变化的能力。高神经可塑性允许个体更好地学习新技能、适应新环境以及从伤害中恢复，这些都是智能的重要组成部分。
- 5. **神经传导速度和效率：**快速且有效的神经信号传导有助于提高信息处理速度和效率，这对解决问题和执行复杂的认知任务至关重要。
- 6. **化学物质平衡：**大脑中的神经递质和其他化学物质的平衡也对维持正常的大脑功能至关重要。例如，多巴胺和血清素等神经递质对于情绪调节、动机和奖励系统非常重要，这些都间接影响着个体的学习能力和创造力。

综上所述，一个生物个体的智商是由大脑的多种特性共同决定的，包括但不限于脑容量、大脑皮层的发达程度、神经连接的复杂性和可塑性、神经传导的速度和效率以及化学物质的平衡状态。然而，值得注意的是，智商只是衡量智力的一个维度，它并不能完全代表一个人的认知能力或成功潜力。此外，环境因素和个人经历也在很大程度上影响着个人的发展和表现。

4. AI记忆与生物记忆对比

4.1 AI记忆与人大脑的差异

该表格有助于系统性地比较人类记忆与AI记忆在存储、处理和演化机制方面的异同，适用于认知科学、人工智能、类脑计算等相关领域的研究与教学。

类型	维度	人类大脑	AI系统
共同点	存储层级划分	分为感觉记忆、短期记忆、长期记忆	分为参数化记忆、上下文记忆、结构化/非结构化记忆
	信息处理逻辑	通过海马体等脑区进行记忆巩固	通过特定算法将新信息转化为长期记忆
	遗忘机制	存在自然遗忘现象	主动删除冗余或过时的记忆
差异点	存储本质	基于神经元及其突触连接等生物结构	基于数字模型参数和数据结构
	学习方式	依赖生活体验、感知、学习等自然过程	依赖数据训练和算法优化
	灵活性与适应性	能灵活适应复杂多变的环境和情境	适应性受限于预设程序、训练数据和算法设计

4.1.1 底层逻辑差异性

操作维度	人类生物记忆	AI记忆系统
物质基础	神经元突触、神经递质、脑区网络	数字模型参数、数据结构、算法逻辑
驱动机制	生物电信号、化学递质动态平衡	预设程序、数学运算、数据输入驱动
灵活性	联想推理（跨模态关联、创造性提取）	严格依赖匹配规则（如向量相似度）
自主性	受意识、情绪、生理状态主动调控	无自主意识，完全依赖外部指令
遗忘逻辑	非线性（受记忆情感价值、使用频率影响）	规则化（时间戳、重要性评分等量化指标）
多模态处理	自然融合感官输入（如嗅觉触发情景记忆）	需单独处理跨模态特征（如多模态编码）

4.2 记忆表示的差异性

记忆类型	人类生物记忆示例及描述	AI记忆系统示例及描述
参数化记忆	本能记忆/先天知识如新生儿吮吸反射、基本生存本能。存储形式：基因编码的神经网络连接（如脊髓反射弧）。	预训练模型参数如GPT-3存储的语言规律、视觉模型的边缘检测能力。存储形式：模型权重矩阵（数字参数）。
上下文结构化记忆	语义记忆/情景记忆如事实性知识（巴黎是法国首都）、个人经历记忆。存储形式：大脑皮层的分布式神经网络（如海马体关联皮层）。	向量数据库/知识图谱如对话历史记录、知识库中的实体关系。存储形式：结构化数据（如JSON、图数据库节点）。
上下文非结构化记忆	感觉记忆/工作记忆如刚看到的图像暂存、短时语音信息。存储形式：短暂神经活动（如视觉皮层的持续放电）。	原始输入数据（未编码）如未处理的文本prompt、原始图像像素值。存储形式：二进制数据（如文本字符串、图像矩阵）。

4.3 六大对应操作

以下是将你提供的“人类生物记忆的六大对应操作”信息整理成结构清晰的表格，便于对比理解人类记忆机制与 AI 记忆操作之间的对应关系：

人类生物记忆操作	对应 AI 操作	机制描述	主要特点
编码（Encoding）	巩固（Consolidation）	通过感官系统将外界信息转化为神经信号（如视觉皮层编码图像）	受注意力、情绪影响（如杏仁核增强情绪记忆编码）
存储（Storage）	更新 + 压缩（Updating + Compression）	短期记忆（海马体）→长期记忆（大脑皮层）的分级存储	存在记忆重组现象（每次提取后可能被修改）
检索（Retrieval）	检索（Retrieval）	通过线索激活神经网络（如情景记忆的时空线索触发）	存在记忆扭曲（如错误归因、暗示性记忆）
遗忘（Forgetting）	遗忘（Forgetting）	突触连接弱化、记忆干扰（前摄/倒摄抑制）	非线性遗忘曲线（符合艾宾浩斯遗忘规律）
重组（Reorganization）	索引（Indexing）	通过语义关联构建记忆网络（如概念层级结构）	睡眠时记忆重组效率提升（海马体与皮层交互）

调控 (Regulation)	压缩 (Compression)	通过表观遗传修饰调节记忆强度（如 CREB 蛋白调控）	情绪调节记忆清晰度（如创伤记忆的强化存储）
--------------------	---------------------	-----------------------------	-----------------------

4.4 典型案例对比

- 巩固操作
 - 人类：学习后睡眠促进记忆固化（海马体向皮层转移）
 - AI：通过梯度下降更新参数实现知识巩固（如微调预训练模型）
- 检索操作
 - 人类："舌尖现象"（知道信息存在但无法提取）
 - AI：检索失败即返回空值（取决于索引完整性）
- 遗忘操作
 - 人类：创伤记忆可能因PTSD难以遗忘（情绪调控异常）
 - AI：按预设策略删除记忆（如FIFO队列淘汰旧数据）

4.5 交叉研究应用

- AI借鉴生物机制
 - 模仿睡眠记忆重组设计"离线训练"算法（如DeepMind的Dreaming训练）
 - 基于遗忘曲线优化AI知识更新策略（如间隔重复学习算法）
- 生物研究借助AI模型
 - 用六大操作框架模拟阿尔茨海默病的记忆衰退模式
 - 通过AI检索算法验证人类记忆提取的神经网络模型

5. Reference

Memory Layers at Scale <https://arxiv.org/abs/2412.09764>

AWESOME: GPU Memory-constrained Long Document Summarization using Memory Mechanism and Global Salient Content <https://arxiv.org/abs/2305.14806>

Keep Me Updated! Memory Management in Long-term Conversations
<https://arxiv.org/abs/2210.08750>

Open Problems in Machine Unlearning for AI Safety <https://arxiv.org/abs/2501.04952>

Memory Layers at Scale <https://arxiv.org/abs/2412.09764>

Judicial knowledge-enhanced magnitude-aware reasoning for numerical legal judgment prediction <https://dl.acm.org/doi/10.1007/s10506-022-09337-4>

Retaining Key Information under High Compression Ratios: Query-Guided Compressor for LLMs
<https://arxiv.org/abs/2406.02376>

Unlearn What You Want to Forget: Efficient Unlearning for LLMs <https://arxiv.org/abs/2310.20150>

Improving Factuality with Explicit Working Memory <https://arxiv.org/abs/2412.18069>

Compress to Impress: Unleashing the Potential of Compressive Memory in Real-World Long-Term Conversations <https://arxiv.org/abs/2402.11975>

Renze Chen, Zhuofeng Wang, Beiquan Cao, Tong Wu, Size Zheng, Xiuhong Li, Xuechao Wei, Shengen Yan, Meng Li, and Yun Liang. 2024c. Arkvale: Efficient generative llm inference with recallable key-value eviction. In *Advances in Neural Information Processing Systems*, volume 37, pages 113134–113155. Curran Associates, Inc.

Wenhu Chen, Zhihao He, Yu Su, Yunyao Yu, William Wang, and Xifeng Yan. 2021b. Hybridqa: A dataset of multi-hop question answering over tabular and textual data. In *Proceedings of the International Conference on Learning Representations (ICLR)*.

Yilong Chen, Guoxia Wang, Junyuan Shang, Shiyao Cui, Zhenyu Zhang, Tingwen Liu, Shuohuan Wang, Yu Sun, Dianhai Yu, and Hua Wu. 2024d. NACL: A general and effective KV cache eviction framework for LLM at inference time. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7913–7926, Bangkok, Thailand. Association for Computational Linguistics.

信息压缩

Xin Cheng, Xun Wang, Xingxing Zhang, Tao Ge, Si-Qing Chen, Furu Wei, Huishuai Zhang, and Dongyan Zhao. 2024. xrag: Extreme context compression for retrieval-augmented generation

with one token. *arXiv preprint arXiv:2405.13792*.

Alexis Chevalier, Alexander Wettig, Anirudh Ajith, and Danqi Chen. 2023. Adapting language models to compress contexts. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 3829–3846, Singapore. Association for Computational Linguistics.

Yu-Neng Chuang, Tianwei Xing, Chia-Yuan Chang, Zirui Liu, Xun Chen, and Xia Hu. 2024. Learning to compress prompt in natural language formats. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 7756–7767, Mexico City, Mexico. Association for Computational Linguistics.

Tsz Ting Chung, Leyang Cui, Lemao Liu, Xinting Huang, Shuming Shi, and Dit-Yan Yeung. 2024. Selection-p: Self-supervised task-agnostic prompt compression for faithfulness and transferability. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 11057–11070, Miami, Florida, USA. Association for Computational Linguistics.

Coze. 2024. Coze: Build your own ai agent. <https://www.coze.cn/>. Accessed: April 19, 2025.

Bhavana Dalvi Mishra, Oyvind Tafjord, and Peter Clark. 2022. Towards teachable reasoning systems: Using a dynamic memory of user feedback for continual system improvement. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 9465–9480, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.

Nelson Cowan. 2001. The magical number 4 in short-term memory: A reconsideration of mental storage capacity. *Behavioral and Brain Sciences*, 24(1):87–114.

Payel Das, Subhajit Chaudhury, Elliot Nelson, Igor Melnyk, Sarathkrishna Swaminathan, Sihui Dai, Aurélie C. Lozano, Georgios Kollias, Vijil Chenthamarakshan, Jirí Navrátil, Soham Dan, and Pin-Yu Chen.

2024. Larimar: Large language models with episodic memory control. In *ICML*

Harry Dong, Xinyu Yang, Zhenyu Zhang, Zhangyang Wang, Yuejie Chi, and Beidi Chen. 2024. Get more with LESS: Synthesizing recurrence with KV cache compression for efficient LLM inference. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 11437–11452. PMLR.

生物学

Ronald L Davis and Yi Zhong. 2017. The biology of forgetting—a perspective. *Neuron*, 95(3):490–503.

Danny Driess, Fei Xia, Mehdi SM Sajjadi, Corey Lynch, Aakanksha Chowdhery, Brian Ichter, Ayzaan Wahid, Jonathan Tompson, Quan Vuong, Tianhe Yu, et al. 2023. Palm-e: An embodied multimodal language model. *arXiv preprint arXiv:2303.03378*.

LLM

N De Cao, W Aziz, and I Titov. 2021. Editing factual knowledge in language models. In *EMNLP 2021-2021 Conference on Empirical Methods in Natural Language Processing, Proceedings*, pages 6491–6506.

Cyprien de Masson D’Autume, Sebastian Ruder, Ling-peng Kong, and Dani Yogatama. 2019. Episodic memory in lifelong language learning. *Advances in Neural Information Processing Systems*, 32.

Alessio Devoto, Yu Zhao, Simone Scardapane, and Pasquale Minervini. 2024. A simple and effective l₂ norm-based strategy for KV cache compression. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 18476–18499, Miami, Florida, USA. Association for Computational Linguistics.

Xuanwen Ding, Jie Zhou, Liang Dou, Qin Chen, Yuan-bin Wu, Arlene Chen, and Liang He. 2024a. Boosting large language models with continual learning for aspect-based sentiment analysis. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 4367–4377, Miami, Florida, USA. Association for Computational Linguistics.

Yiran Ding, Li Lyna Zhang, Chengruidong Zhang, Yuanyuan Xu, Ning Shang, Jiahang Xu, Fan Yang, and Mao Yang. 2024b. Longrope: Extending llm context window beyond 2 million tokens.

Yiming Du, Bingbing Wang, Yang He, Bin Liang, Baojun Wang, Zhongyang Li, Lin Gui, Jeff Z. Pan, Ruifeng Xu, and Kam-Fai Wong. 2025. Bridging the long-term gap: A memory-active policy for multi-session task-oriented dialogue. *arXiv preprint arXiv:2505.20231*.

Yiming Du, Hongru Wang, Zhengyi Zhao, Bin Liang, Baojun Wang, Wanjun Zhong, Zezhong Wang, and Kam-Fai Wong. 2024. Perltqa: A personal long-term memory dataset for memory classification, retrieval, and synthesis in question answering. *arXiv preprint arXiv:2402.16288*.

Yadin Dudai, Avi Karni, and Jan Born. 2015. The consolidation and transformation of memory. *Neuron*, 88(1):20–32.

Junfeng Fang, Houcheng Jiang, Kun Wang, Yunshan Ma, Jie Shi, Xiang Wang, Xiangnan He, and Tat-Seng Chua. 2025. Alphaedit: Null-space constrained model editing for language models. In *The Thirteenth International Conference on Learning Representations*.

Yujie Feng, Xu Chu, Yongxin Xu, Guangyuan Shi, Bo Liu, and Xiao-Ming Wu. 2024. TaSL: Continual dialog state tracking via task skill localization and consolidation. In *Proceedings of the 62nd*

Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), pages 1266–1279, Bangkok, Thailand. Association for Computational Linguistics.

Nicholas T Franklin, Kenneth A Norman, Charan Ranganath, Jeffrey M Zacks, and Samuel J Gershman. 2020. Structured event memory: A neuro-symbolic model of event cognition. *Psychological Review*, 127(3):327–361.

u Fu, Zefan Cai, Abedelkadir Asi, Wayne Xiong, Yue Dong, and Wen Xiao. 2025. Not all heads matter: A head-level KV cache compression method with integrated retrieval and reasoning. In *The Thirteenth International Conference on Learning Representations*.

Suyu Ge, Yunan Zhang, Liyuan Liu, Minjia Zhang, Jiawei Han, and Jianfeng Gao. 2024. Model tells you what to discard: Adaptive KV cache compression for LLMs. In *The Twelfth International Conference on Learning Representations*.

Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. 2024. Hipporag: Neurobiologically inspired long-term memory for large language models. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*.

Bernal Jiménez Gutiérrez, Yiheng Shu, Weijian Qi, Sizhe Zhou, and Yu Su. 2025. From rag to memory: Non-parametric continual learning for large language models. *arXiv preprint arXiv:2502.14802*.

Kaiqiao Han, Tianqing Fang, Zhaowei Wang, Yangqiu Song, and Mark Steedman. 2025. Concept-reversed Winograd schema challenge: Evaluating and improving robust reasoning in large language models via abstraction. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pages 229–243, Albuquerque, New Mexico. Association for Computational Linguistics.

Yongchang Hao, Mengyao Zhai, Hossein Hajimirsadeghi, Sepidehsadat Hosseini, and Frederick Tung. 2025. Radar: Fast long-context decoding for any transformer. In *The Thirteenth International Conference on Learning Representations*.

Zihong He, Weizhe Lin, Hao Zheng, Fan Zhang, Matt W. Jones, Laurence Aitchison, Xuhai Xu, Miao Liu, Per Ola Kristensson, and Junxiao Shen. 2024d. Human-inspired perspectives: A survey on ai long-term memory. *arXiv preprint arXiv:2411.00489*.

Yuki Hou, Haruki Tamoto, and Homei Miyashita. 2024.

"my agent understands me better": Integrating dynamic human-like memory recall and consolidation in llm-based agents. In *Extended Abstracts of the CHI Conference on Human Factors in Computing Systems*, pages 1–7. ACM.

Zhijian Hou, Lei Ji, Difei Gao, Wanjun Zhong, Kun Yan, Chao Li, Wing-Kwong Chan, Chong-Wah Ngo, Nan Duan, and Mike Zheng Shou. 2023. Groundnlq@ ego4d natural language queries challenge 2023. *arXiv preprint arXiv:2306.15255*.

Chenxu Hu, Jie Fu, Chenzhuang Du, Simian Luo, Junbo Zhao, and Hang Zhao. 2023. Chatdb: Augmenting llms with databases as their symbolic memory.

Luyang Huang, Shuyang Cao, Nikolaus Parulian, Heng Ji, and Lu Wang. 2021. Efficient attentions for long document summarization. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 1419–1436, Online. Association for Computational Linguistics.

Qiushi Huang, Shuai Fu, Xubo Liu, Wenwu Wang, Tom Ko, Yu Zhang, and Lilian Tang. 2023a. Learning retrieval augmentation for personalized dialogue generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2523–2540, Singapore. Association for Computational Linguistics.

Xiusheng Huang, Yequan Wang, Jun Zhao, and Kang Liu. 2024. Commonsense knowledge editing based on free-text in llms. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 14870–14880.

Yunpeng Huang, Jingwei Xu, Junyu Lai, Zixu Jiang, Taolue Chen, Zenan Li, Yuan Yao, Xiaoxing Ma, Li-juan Yang, Hao Chen, et al. 2023b. Advancing transformer architecture in long-context large language models: A comprehensive survey. *arXiv preprint arXiv:2311.12351*.

Kai Tzu iunn Ong, Namyoung Kim, Minju Gwak, Hyungjoo Chae, Taeyoon Kwon, Yohan Jo, Seung won Hwang, Dongha Lee, and Jinyoung Yeo. 2025. Towards lifelong dialogue agents via timeline-based memory management. In *Proceedings of the 2025 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, Mexico City, Mexico. Association for Computational Linguistics.

Gautier Izacard, Mathilde Caron, Lucas Hosseini, Sebastian Riedel, Piotr Bojanowski, Armand Joulin, and Edouard Grave. 2021. Unsupervised dense information retrieval with contrastive learning. *arXiv preprint arXiv:2112.09118*.

Jiabao Ji, Yujian Liu, Yang Zhang, Gaowen Liu, Ramana Kompella, Sijia Liu, and Shiyu Chang. 2024. Reversing the forget-retain objectives: An efficient llm unlearning framework from logit difference. *Advances in Neural Information Processing Systems*, 37:12581–12611.

Jinghan Jia, Jiancheng Liu, Yihua Zhang, Parikshit Ram, Nathalie Baracaldo Angel, and Sijia Liu. 2024a. Wa- gle: Strategic weight attribution for effective and modular unlearning in large language models. In *Annual Conference on Neural Information Processing Systems*.

Huiqiang Jiang, Qianhui Wu, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. 2023a. LLMlingua: Compress- ing prompts for accelerated inference of large lan- guage models. In *Proceedings of the 2023 Confer- ence on Empirical Methods in Natural Language Pro- cessing*, pages 13358–13376, Singapore. Association for Computational Linguistics.

Xun Jiang, Feng Li, Han Zhao, Jiaying Wang, Jun Shao, Shihao Xu, Shu Zhang, Weiling Chen, Xavier Tang, Yize Chen, Mengyue Wu, Weizhi Ma, Mengdi Wang, and Tianqiao Chen. 2024b. Long term memory: The foundation of ai self-evolution. *arXiv preprint arXiv:2410.15665*.

Bowen Jin, Jinsung Yoon, Jiawei Han, and Sercan O Arik. 2025. Long-context LLMs meet RAG: Over- coming challenges for long inputs in RAG. In *The Thirteenth International Conference on Learning Representations*.

Mingyu Jin, Weidi Luo, Sitao Cheng, Xinyi Wang, Wenyue Hua, Ruixiang Tang, William Yang Wang, and Yongfeng Zhang. 2024a. Disentangling mem- ory and reasoning ability in large language models. *arXiv preprint arXiv:2411.13504*.

Zhuoran Jin, Pengfei Cao, Chenhao Wang, Zhitao He, Hongbang Yuan, Jiachun Li, Yubo Chen, Kang Liu, and Jun Zhao. 2024b. RWKU: Benchmarking real- world knowledge unlearning for large language mod- els. In *The Thirty-eight Conference on Neural In- formation Processing Systems Datasets and Bench- marks Track*.

Christopher Kiley and Colleen M Parks. 2022. Mechanisms of memory updating: State dependency vs. reconsolidation. *Journal of cognition*, 5(1):7.

Minsoo Kim, Kyuhong Shim, Jungwook Choi, and Simyung Chang. 2024b. InfiniPot: Infinite context processing on memory-constrained LLMs. In *Pro- ceedings of the 2024 Conference on Empirical Meth- ods in Natural Language Processing*, pages 16046–16060, Miami, Florida, USA. Association for Com- putational Linguistics.

Seo Hyun Kim, Keummin Ka, Yohan Jo, Seung-won Hwang, Dongha Lee, and Jinyoung Yeo. 2024c. Ever- evolving memory by blending and refining the past. *arXiv preprint arXiv:2403.04787*.

kingjulio8238. 2025. Memary. <https://github.com/kingjulio8238/Memary>. 2025-04-17.

James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A Rusu, Kieran Milan, John Quan, Tiago Ramalho, Ag- nieszka Grabska-Barwinska, et al. 2017. Over- coming catastrophic forgetting in neural networks. *Proceedings of the national academy of sciences*, 114(13):3521–3526.

Dharshan Kumaran, Demis Hassabis, and James L Mc- Clelland. 2016. What learning systems do intelligent agents need? complementary learning systems theory updated. *Trends in Cognitive Sciences*, 20(7):512– 534.

Tom Kwiatkowski, Jennimaria Palomaki, Olivia Red- field, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Ken- ton Lee, et al. 2019. Natural questions: a benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453– 466.

Jack Lanchantin, Shubham Toshniwal, Jason Weston, Sainbayar Sukhbaatar, et al. 2023. Learning to reason and memorize with self-notes. *Advances in Neural Information Processing Systems*, 36:11891–11911.

Kuang-Huei Lee, Xinyun Chen, Hiroki Furuta, John Canny, and Ian Fischer. 2024b. A human- inspired reading agent with gist memory of very long contexts. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 26396–26415. PMLR.

Mingcong Lei, Yiming Zhao, Ge Wang, Zhixin Mai, Shuguang Cui, Yatong Han, and Jinke Ren. 2025. Stma: A spatio-temporal memory agent for long- horizon embodied task planning. *arXiv preprint arXiv:2502.10177*.

Yaniv Leviathan, Matan Kalman, and Yossi Matias. 2025. Selective attention improves transformer. In *The Thirteenth International Conference on Learning Representations*.

Hao Li, Chenghao Yang, An Zhang, Yang Deng, Xiang Wang, and Tat-Seng Chua. 2024a. Hello again! llm- powered personalized agent for long-term dialogue. *arXiv preprint arXiv:2406.05925*.

Jiaqi Li, Mengmeng Wang, Zilong Zheng, and Muhan Zhang. 2024b. LooGLE: Can long-context language models understand long contexts? In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16304–16333, Bangkok, Thailand. Association for Computational Linguistics.

Na Li, Chunyi Zhou, Yansong Gao, Hui Chen, Zhi Zhang, Boyu Kuang, and Anmin Fu. 2025. Machine unlearning: Taxonomy, metrics, applications, challenges, and prospects. *IEEE Transactions on Neural Networks and Learning Systems*, pages 1–21.

Machine Unlearning: Taxonomy, Metrics, Applications, Challenges, and Prospects <https://arxiv.org/abs/2403.08254>

Xiaonan Li and Xipeng Qiu. 2023. Mot: Memory-of- thought enables chatgpt to self-improve. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 6354–6374, Singapore. Association for Computational Linguistics.

Xiaopeng Li, Shasha Li, Shezheng Song, Jing Yang, Jun Ma, and Jie Yu. 2024e. Pmet: Precise model editing in a transformer. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 18564–18572.

Yuhong Li, Yingbing Huang, Bowen Yang, Bharat Venkitesh, Acyr Locatelli, Hanchen Ye, Tianle Cai, Patrick Lewis, and Deming Chen. 2024h. Snapkv: Llm knows what you are looking for before generation. In *Advances in Neural Information Processing Systems*, volume 37, pages 22947–22970. Curran Associates, Inc.

Zaijing Li, Yuquan Xie, Rui Shao, Gongwei Chen, Dongmei Jiang, and Liqiang Nie. 2024i. Optimus- 1: Hybrid multimodal memory empowered agents excel in long-horizon tasks. *arXiv preprint arXiv:2408.03615*.

Zhuowan Li, Cheng Li, Mingyang Zhang, Qiaozhu Mei, and Michael Bendersky. 2024k. Retrieval augmented generation or long-context LLMs? a comprehensive study and hybrid approach. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: Industry Track*, pages 881– 893, Miami, Florida, US. Association for Computational Linguistics.

Zongqian Li, Yinhong Liu, Yixuan Su, and Nigel Collier. 2024l. Prompt compression for large language models: A survey.

Yinhong Liu, Zexuan Qiu, Zhongyang Li, Quanyu Dai, Jieming Zhu, Minda Hu, Menglin Yang, and Irwin King. 2025a. A survey of personalized large language models: Progress and future directions. *arXiv preprint arXiv:2502.11528*.

Ziyao Liu, Huanyi Ye, Chen Chen, Yongsen Zheng, and Kwok-Yan Lam. 2025b. Threats, attacks, and defenses in machine unlearning: A survey. *IEEE Open Journal of the Computer Society*, 6:413–425.

Liyuan Wang, Xingxing Zhang, Hang Su, and Jun Zhu. 2024b. A comprehensive survey of continual learning: Theory, method and application. *IEEE Transactions on Pattern Analysis and Machine Intelligence*.

Rongwu Xu, Zehan Qi, Zhijiang Guo, Cunxiang Wang, Hongru Wang, Yue Zhang, and Wei Xu. 2024b. Knowledge conflicts for llms: A survey. *arXiv preprint arXiv:2403.08319*.

Zeyu Zhang, Xiaohe Bo, Chen Ma, Rui Li, Xu Chen, Quanyu Dai, Jieming Zhu, Zhenhua Dong, and Ji-Rong Wen. 2024f. A survey on the memory mechanism of large language model-based agents. *arXiv preprint arXiv:2404.13501*.

AI「记忆」全景解读：三大形态、六大操作 <https://zhuanlan.zhihu.com/p/1906771061212688820>

万字长文总结大模型知识编辑领域最新研究进

<https://blog.csdn.net/tmb8z9vdm66wh68vx1/article/details/138981986>

MEMIT: 批量编辑变压器记忆文档指南

https://blog.csdn.net/gitblog_00940/article/details/142010585

语言模型知识编辑的鲁棒性研究<https://www.51cto.com/aigc/46.html>

GPT-4 Technical Report <https://arxiv.org/abs/2303.08774>

Automatic Chain of Thought Prompting in Large Language Models

<https://arxiv.org/abs/2210.03493>

麦肯锡是如何沟通的？ http://www.360doc.com/content/23/0524/18/31263000_1081978090.shtml